

Structure from Motion

Outline

- Bundle Adjustment
- Ambiguities in Reconstruction
- Affine Factorization
- Extensions

Structure from motion

Recover both 3D scene geometry **and** camera positions




SLAM: Simultaneous Localization and Mapping




Sequence 00

Recall: 2-view stereo




An annoying detail



What to do when ray's don't intersect?

Minimize reprojection error




$$\min_{\mathbf{X}} f(\mathbf{X}) = ||\mathbf{x}_1 - Proj(\mathbf{X}, M_1)||^2 + ||\mathbf{x}_2 - Proj(\mathbf{X}, M_2)||^2$$

Perspective projection equations where $M_1 = 3 \times 4$ matrix of extrinsics (R_1, t_1) and intrinsics (f_1) of camera 1

For this equation, its easier to parameterize camera position in absolute coordinates instead of relative ones

Minimize reprojection error




$$\min_{\mathbf{X}} f(\mathbf{X}) = ||\mathbf{x}_1 - Proj(\mathbf{X}, M_1)||^2 + ||\mathbf{x}_2 - Proj(\mathbf{X}, M_2)||^2$$

Perspective projection equations where $M_1 = 3 \times 4$ matrix of extrinsics (R_1, t_1) and intrinsics (f_1) of camera 1

For this equation, its easier to parameterize camera position in absolute coordinates instead of relative ones

Generalize triangulation to multiple points from multiple cameras




$$\min_{\mathbf{X}_1, \mathbf{X}_2, \dots} \sum_{i=1}^m \sum_{j=1}^n ||\mathbf{x}_{ij} - Proj(\mathbf{X}_j, M_i)||^2$$

As written, could this minimization be done independently for each 3D point?


Bundle adjustment

Minimize reprojection error over multiple 3D points and **cameras**



$$\min_{\mathbf{X}_1, \mathbf{X}_2, \dots, M_1, M_2, \dots} \sum_{i=1}^m \sum_{j=1}^n \|\mathbf{x}_{ij} - Proj(\mathbf{X}_j, M_i)\|^2$$


Basic SFM pipeline



1. Find candidate correspondences (interest points + descriptor matches)
2. Select subset that are consistent with epipolar constraints on pairs of images (RANSAC + fundamental matrix)
3. Solve for 3D points and camera that minimize reprojection error

(Lots of variants; e.g., iteratively build map of 3D points and cameras as new images arrive)


$$\min_{\mathbf{X}_1, \mathbf{X}_2, \dots, M_1, M_2, \dots} \sum_{i=1}^m \sum_{j=1}^n ||\mathbf{x}_{ij} - Proj(\mathbf{X}_j, M_i)||^2$$



Bundle adjustment: nonlinear least-squares

Encode visibility graphs of what points are seen in what images, i.e.

- Features: A,B,C,D,E
- Images: 1,2,3



Implies jacobian of error function is sparse

	A	B	C	D	E	1	2	3	4
A1	■					■			
A2	■						■		
B1		■				■			
B2		■					■		
B4		■							■
C1			■			■			
C3			■					■	
D2				■			■		
D3				■				■	
D4				■					■
E3					■			■	
E4					■				■

Excellent reference

Bundle Adjustment — A Modern Synthesis

Bill Triggs¹, Philip McLauchlan², Richard Hartley³ and Andrew Fitzgibbon⁴

¹ INRIA Rhône-Alpes, 655 avenue de l'Europe, 38330 Montbonnot, France.

Bill.Triggs@inrialpes.fr \diamond <http://www.inrialpes.fr/movi/people/Triggs>

² School of Electrical Engineering, Information Technology & Mathematics
University of Surrey, Guildford, GU2 5XH, U.K.

P.McLauchlan@ee.surrey.ac.uk \diamond <http://www.ee.surrey.ac.uk/Personal/P.McLauchlan>

³ General Electric CRD, Schenectady, NY, 12301

hartley@crd.ge.com

⁴ Dept of Engineering Science, University of Oxford, 19 Parks Road, OX1 3PJ, U.K.

awf@robots.ox.ac.uk \diamond <http://www.robots.ox.ac.uk/awf>

Outline

- Bundle Adjustment
- Ambiguities in Reconstruction
- Affine Factorization
- Extensions

Recall camera projection

$$\lambda \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} fs_x & fs_\theta & o_x \\ 0 & fs_y & o_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

$$= K_{3 \times 3} \begin{bmatrix} R_{3 \times 3} & T_{3 \times 1} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

$$= M_{3 \times 4} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

$$\mathbf{x} \equiv M\mathbf{X}$$

$$x = \frac{m_1^T X}{m_3^T X}$$

$$y = \frac{m_2^T X}{m_3^T X}$$

Structure from motion ambiguity

- If we scale the entire scene by some factor k and, at the same time, scale the camera matrices by the factor of $1/k$, the 2D homogenous vector remains exactly the same:

$$\mathbf{x} \equiv M\mathbf{X}$$
$$M\mathbf{X} = \left(\frac{1}{k}M\right)(k\mathbf{X})$$

It is impossible to recover the absolute scale of the scene!



Structure from motion ambiguity

- If we scale the entire scene by some factor k and, at the same time, scale the camera matrices by the factor of $1/k$, the 2D homogenous vector remains exactly the same:
- More generally: if we transform the scene using a transformation \mathbf{Q} and apply the inverse transformation to the camera matrices, nothing changes

$$\mathbf{x} \equiv M\mathbf{X}$$

$$M\mathbf{X} = (MQ^{-1})(Q\mathbf{X})$$



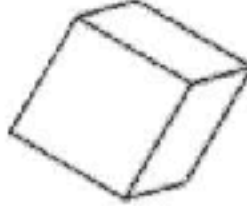
Recall: transformations in 3D

Similarity 7dof	$\begin{bmatrix} sR & \mathbf{t} \\ 0^T & 1 \end{bmatrix}$		Preserves angles, ratios of length
Euclidean 6dof	$\begin{bmatrix} R & \mathbf{t} \\ 0^T & 1 \end{bmatrix}$		Preserves angles, lengths




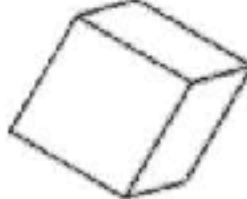
Make use of 3D homogenous coordinates

$$\begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Recall: transformations in 3D

Affine 12dof	$\begin{bmatrix} A & \mathbf{t} \\ 0^T & 1 \end{bmatrix}$		Preserves parallelism, volume ratios
Similarity 7dof	$\begin{bmatrix} sR & \mathbf{t} \\ 0^T & 1 \end{bmatrix}$		Preserves angles, ratios of length
Euclidean 6dof	$\begin{bmatrix} R & \mathbf{t} \\ 0^T & 1 \end{bmatrix}$		Preserves angles, lengths

Recall: transformations in 3D

Projective 15dof	$\begin{bmatrix} A & t \\ v & v \end{bmatrix}$		Preserves intersection and tangency
Affine 12dof	$\begin{bmatrix} A & t \\ 0^T & 1 \end{bmatrix}$		Preserves parallelism, volume ratios
Similarity 7dof	$\begin{bmatrix} sR & t \\ 0^T & 1 \end{bmatrix}$		Preserves angles, ratios of length
Euclidean 6dof	$\begin{bmatrix} R & t \\ 0^T & 1 \end{bmatrix}$		Preserves angles, lengths

Normalize by last coordinate to recover 3D points


$$\begin{bmatrix} \lambda x \\ \lambda y \\ \lambda z \\ \lambda \end{bmatrix}$$

Back to ambiguities for 3D reconstruction

$$\mathbf{x} = \mathbf{P}\mathbf{X} = (\mathbf{P}\mathbf{Q}^{-1})\mathbf{Q}\mathbf{X}$$

Projective


$$\begin{bmatrix} \mathbf{A} & \mathbf{t} \\ \mathbf{v}^\top & v \end{bmatrix}$$



Preserves intersection

Affine

$$\begin{bmatrix} \mathbf{A} & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix}$$



Preserves parallelism,

Similarity


$$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix}$$



Preserves angles, ratios

Euclidean

$$\begin{bmatrix} \mathbf{R} & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix}$$




Preserves angles.

- With no constraints on the camera calibration matrix or on the scene, we get a *projective* reconstruction
- Need additional information to *upgrade* the reconstruction to affine, similarity, or Euclidean


Projective ambiguity

$$x \equiv M\mathbf{X} = (MQ^{-1})(Q\mathbf{X})$$



$$\mathbf{Q}_p = \begin{bmatrix} \mathbf{A} & \mathbf{t} \\ \mathbf{v}^\top & v \end{bmatrix}$$


Projective ambiguity



(straight line are preserved)



Affine ambiguity

$$x \equiv M\mathbf{X} = (MQ^{-1})(Q\mathbf{X})$$



$$\mathbf{Q}_A = \begin{bmatrix} \mathbf{A} & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix}$$


Affine ambiguity



(parallel lines are preserved)




Similarity ambiguity

$$x \equiv M\mathbf{X} = (MQ^{-1})(Q\mathbf{X})$$



$$\mathbf{Q}_s = \begin{bmatrix} s\mathbf{R} & \mathbf{t} \\ \mathbf{0}^\top & 1 \end{bmatrix}$$

Similarity ambiguity





angles+lengths preserved (but can't recover world coordinate system)

Outline

- Bundle Adjustment
- Ambiguities in Reconstruction
- Affine Factorization
- Large-scale SFM


Structure from motion

- Let's start with *affine cameras* (the math is easier)



Recall: Orthographic Projection

Special case of perspective projection



- Another common notation for a projection matrix:

$$\begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & & & \\ & 1 & & \\ & & & \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \Rightarrow (x, y)$$

Sometimes notationally convenient because 2D homogenous coordinates allow us to write 2D transformations as matrix multiplication

Recall: affine cameras

Model as 3D affine transformation + orthographic projection + 2D affine transformation

$$\begin{aligned} \begin{bmatrix} x \\ y \end{bmatrix} &= \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} 1 & & & \\ & 1 & & \\ & & & 1 \end{bmatrix} \begin{bmatrix} \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \\ &= \begin{bmatrix} a_{11} & a_{12} & a_{13} & b_1 \\ a_{21} & a_{22} & a_{23} & b_2 \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \end{aligned}$$

Affine camera defined by 8 parameters

Affine cameras

Model as 3D affine transformation + orthographic projection + 2D affine transformation

$$\begin{aligned}
 \begin{bmatrix} x \\ y \end{bmatrix} &= \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} 1 & & & \\ & 1 & & \\ & & & 1 \end{bmatrix} \begin{bmatrix} \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \\
 &= \begin{bmatrix} a_{11} & a_{12} & a_{13} & b_1 \\ a_{21} & a_{22} & a_{23} & b_2 \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \\
 &= \begin{bmatrix} a_{11} & a_{12} & a_{13} & b_1 \\ a_{21} & a_{22} & a_{23} & b_2 \\ & & & 1 \end{bmatrix} Q_a Q_a^{-1} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}, \quad Q_a = \begin{bmatrix} \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ & & & 1 \end{bmatrix}
 \end{aligned}$$

Q_a must be 3D affine transformation (12 parameters) in order to maintain structure of affine projection matrix

Affine cameras

Model as 3D affine transformation + orthographic projection + 2D affine transformation

$$\begin{aligned}
 \begin{bmatrix} x \\ y \end{bmatrix} &= \begin{bmatrix} \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} 1 & & & \\ & 1 & & \\ & & & 1 \end{bmatrix} \begin{bmatrix} \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \\
 &= \begin{bmatrix} a_{11} & a_{12} & a_{13} & b_1 \\ a_{21} & a_{22} & a_{23} & b_2 \\ & & & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \\
 &= \begin{bmatrix} a_{11} & a_{12} & a_{13} & b_1 \\ a_{21} & a_{22} & a_{23} & b_2 \\ & & & 1 \end{bmatrix} Q_a Q_a^{-1} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}, \quad Q_a = \begin{bmatrix} \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot \\ & & & 1 \end{bmatrix} \\
 &= \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \end{bmatrix}
 \end{aligned}$$

$$\mathbf{x} = \mathbf{A}\mathbf{X} + \mathbf{b}$$


2D points = linear transformation of 3D points + 2D translation

Affine structure from motion (my *favorite* vision algorithm)

- Given: m images of n fixed 3D points:

$$\mathbf{x}_{ij} = \mathbf{A}_i \mathbf{X}_j + \mathbf{b}_i, \quad i = 1, \dots, m, \quad j = 1, \dots, n$$

- Problem: use the mn correspondences \mathbf{x}_{ij} to estimate m projection matrices \mathbf{A}_i and translation vectors \mathbf{b}_i , and n points \mathbf{X}_j



of points
→

of
images
↓

$$\begin{bmatrix} \mathbf{x}_{11} & \mathbf{x}_{12} & \dots \\ \mathbf{x}_{21} & \mathbf{x}_{22} & \dots \\ \vdots & \vdots & \ddots \end{bmatrix}$$

Affine structure from motion

- Given: m images of n fixed 3D points:

$$\mathbf{x}_{ij} = \mathbf{A}_i \mathbf{X}_j + \mathbf{b}_i, \quad i = 1, \dots, m, \quad j = 1, \dots, n$$

- Problem: use the mn correspondences \mathbf{x}_{ij} to estimate m projection matrices \mathbf{A}_i and translation vectors \mathbf{b}_i , and n points \mathbf{X}_j
- The reconstruction is defined up to an arbitrary 3D *affine* transformation \mathbf{Q} (12 degrees of freedom):

$$\begin{bmatrix} \mathbf{A} & \mathbf{b} \\ \mathbf{0} & \mathbf{1} \end{bmatrix} \rightarrow \begin{bmatrix} \mathbf{A} & \mathbf{b} \\ \mathbf{0} & \mathbf{1} \end{bmatrix} \mathbf{Q}^{-1}, \quad \begin{pmatrix} \mathbf{X} \\ \mathbf{1} \end{pmatrix} \rightarrow \mathbf{Q} \begin{pmatrix} \mathbf{X} \\ \mathbf{1} \end{pmatrix}$$

- How many knowns are there?
- How many unknowns?

Affine structure from motion

- Given: m images of n fixed 3D points:

$$\mathbf{x}_{ij} = \mathbf{A}_i \mathbf{X}_j + \mathbf{b}_i, \quad i = 1, \dots, m, \quad j = 1, \dots, n$$

- Problem: use the mn correspondences \mathbf{x}_{ij} to estimate m projection matrices \mathbf{A}_i and translation vectors \mathbf{b}_i , and n points \mathbf{X}_j
- The reconstruction is defined up to an arbitrary 3D *affine* transformation \mathbf{Q} (12 degrees of freedom):

$$\begin{bmatrix} \mathbf{A} & \mathbf{b} \\ \mathbf{0} & \mathbf{1} \end{bmatrix} \rightarrow \begin{bmatrix} \mathbf{A} & \mathbf{b} \\ \mathbf{0} & \mathbf{1} \end{bmatrix} \mathbf{Q}^{-1}, \quad \begin{pmatrix} \mathbf{X} \\ \mathbf{1} \end{pmatrix} \rightarrow \mathbf{Q} \begin{pmatrix} \mathbf{X} \\ \mathbf{1} \end{pmatrix}$$

- We have $2mn$ knowns and $8m + 3n$ unknowns (minus 12 dof for affine ambiguity)
- Thus, we must have $2mn \geq 8m + 3n - 12$
- For $m=2$ views, we need $n=4$ point correspondences


Affine structure from motion

Let's try centering the points in each 2D image

of points
→

of images
↓

$$\begin{bmatrix} \mathbf{x}_{11} & \mathbf{x}_{12} & \dots \\ \mathbf{x}_{21} & \mathbf{x}_{22} & \dots \\ \vdots & \vdots & \ddots \end{bmatrix}$$



$$\hat{\mathbf{x}}_{ij} = \mathbf{x}_{ij} - \frac{1}{n} \sum_k^n \mathbf{x}_{ik}$$

Plug the following into the above expression: $\mathbf{x}_{ij} = \mathbf{A}_i \mathbf{X}_j + \mathbf{b}_i$

[on board]

Affine structure from motion

- Centering: subtract the centroid of the image points

$$\begin{aligned}\hat{\mathbf{x}}_{ij} &= \mathbf{x}_{ij} - \frac{1}{n} \sum_{k=1}^n \mathbf{x}_{ik} = \mathbf{A}_i \mathbf{X}_j + \mathbf{b}_i - \frac{1}{n} \sum_{k=1}^n (\mathbf{A}_i \mathbf{X}_k + \mathbf{b}_i) \\ &= \mathbf{A}_i \left(\mathbf{X}_j - \frac{1}{n} \sum_{k=1}^n \mathbf{X}_k \right) = \mathbf{A}_i \hat{\mathbf{X}}_j\end{aligned}$$

- After centering, each normalized point \mathbf{x}_{ij} is related to the 3D point \mathbf{X}_j by


$$\hat{\mathbf{x}}_{ij} = \mathbf{A}_i \hat{\mathbf{X}}_j$$


Given a set of 2D correspondences, simply center them in each image
Affine image projection now becomes linear (represent \mathbf{A}_i by a 2x3 matrix and $\hat{\mathbf{x}}_{ij}$ is 2 vector)

Affine structure from motion

- Let's create a $2m \times n$ data (measurement) matrix:

$$\mathbf{D} = \begin{bmatrix} \hat{\mathbf{x}}_{11} & \hat{\mathbf{x}}_{12} & \cdots & \hat{\mathbf{x}}_{1n} \\ \hat{\mathbf{x}}_{21} & \hat{\mathbf{x}}_{22} & \cdots & \hat{\mathbf{x}}_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ \hat{\mathbf{x}}_{m1} & \hat{\mathbf{x}}_{m2} & \cdots & \hat{\mathbf{x}}_{mn} \end{bmatrix}$$


points (n)


cameras ($2m$)

Affine structure from motion

- Let's create a $2m \times n$ data (measurement) matrix:


$$\mathbf{D} = \begin{bmatrix} \hat{\mathbf{x}}_{11} & \hat{\mathbf{x}}_{12} & \cdots & \hat{\mathbf{x}}_{1n} \\ \hat{\mathbf{x}}_{21} & \hat{\mathbf{x}}_{22} & \cdots & \hat{\mathbf{x}}_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ \hat{\mathbf{x}}_{m1} & \hat{\mathbf{x}}_{m2} & \cdots & \hat{\mathbf{x}}_{mn} \end{bmatrix} = \begin{bmatrix} \mathbf{A}_1 \\ \mathbf{A}_2 \\ \vdots \\ \mathbf{A}_m \end{bmatrix} \begin{bmatrix} \mathbf{X}_1 & \mathbf{X}_2 & \cdots & \mathbf{X}_n \end{bmatrix}$$


cameras
($2m \times 3$)

points ($3 \times n$)


The measurement matrix $\mathbf{D} = \mathbf{MS}$ must have rank 3!


Fundamental Decomposition






To reduce to rank 3, we just need to set all the singular values to 0 except for the first 3





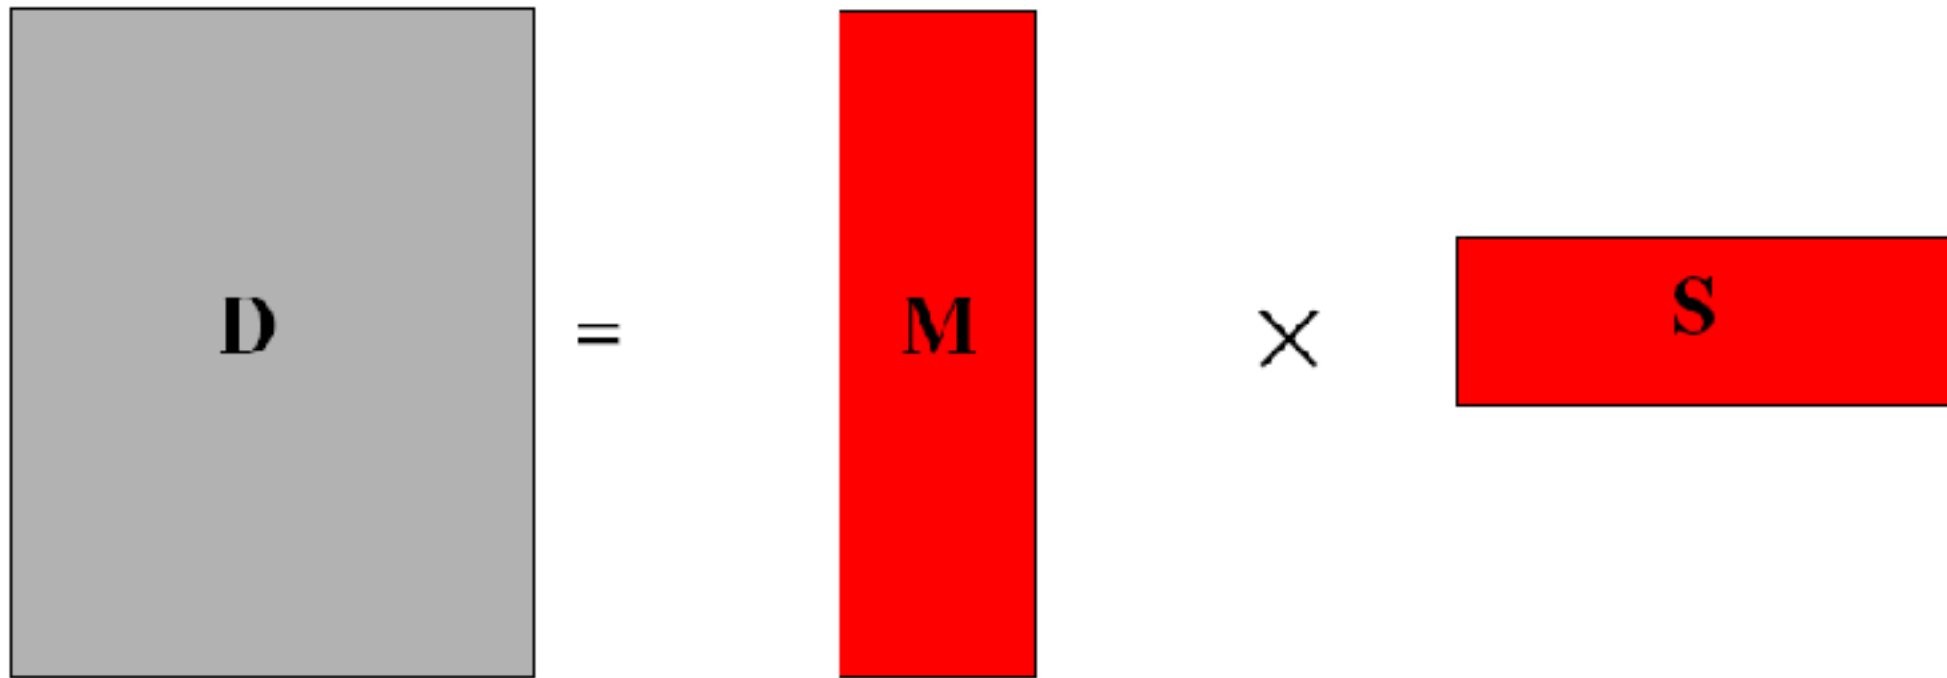
Possible decomposition:

$$M = U_3, \quad S = \Sigma_3 V_3^T$$



This decomposition minimizes $\|\mathbf{D} - \mathbf{MS}\|^2$

Affine ambiguity




The diagram shows a gray square labeled **D** on the left. To its right is an equals sign. Further right is a red vertical rectangle labeled **M**. To the right of **M** is a multiplication symbol \times . To the right of the multiplication symbol is a red horizontal rectangle labeled **S**. This visualizes the equation $D = M \times S$.

- The decomposition is not unique. We get the same **D** by using any 3×3 matrix **C** and applying the transformations $\mathbf{M} \rightarrow \mathbf{MC}$, $\mathbf{S} \rightarrow \mathbf{C}^{-1}\mathbf{S}$
- That is because we have only an affine transformation and we have not enforced any Euclidean constraints (like forcing the image axes to be perpendicular, for example)

Eliminating the affine ambiguity

enforce image axes to be orthonormal and length 1



$$\mathbf{a}_1 \cdot \mathbf{a}_2 = 0$$

$$|\mathbf{a}_1|^2 = |\mathbf{a}_2|^2 = 1$$

$$M = \begin{bmatrix} A_1 \\ A_2 \\ \vdots \end{bmatrix}$$


We want C such that $A_i C C^T A_i^T = Id$

Write out optimization as min

$$\min_L ||A_i L A_i^T - Id||^2$$

Eliminating the affine ambiguity

- Orthographic: image axes are perpendicular and scale is 1




- This translates into $3m$ equations in $\mathbf{L} = \mathbf{C}\mathbf{C}^T$:
$$\mathbf{A}_i \mathbf{L} \mathbf{A}_i^T = \mathbf{Id}, \quad i = 1, \dots, m$$

- Solve for \mathbf{L}
- Recover \mathbf{C} from \mathbf{L} by Cholesky decomposition: $\mathbf{L} = \mathbf{C}\mathbf{C}^T$ (in practice, easy to do)
- Update \mathbf{M} and \mathbf{S} : $\mathbf{M} = \mathbf{M}\mathbf{C}$, $\mathbf{S} = \mathbf{C}^{-1}\mathbf{S}$

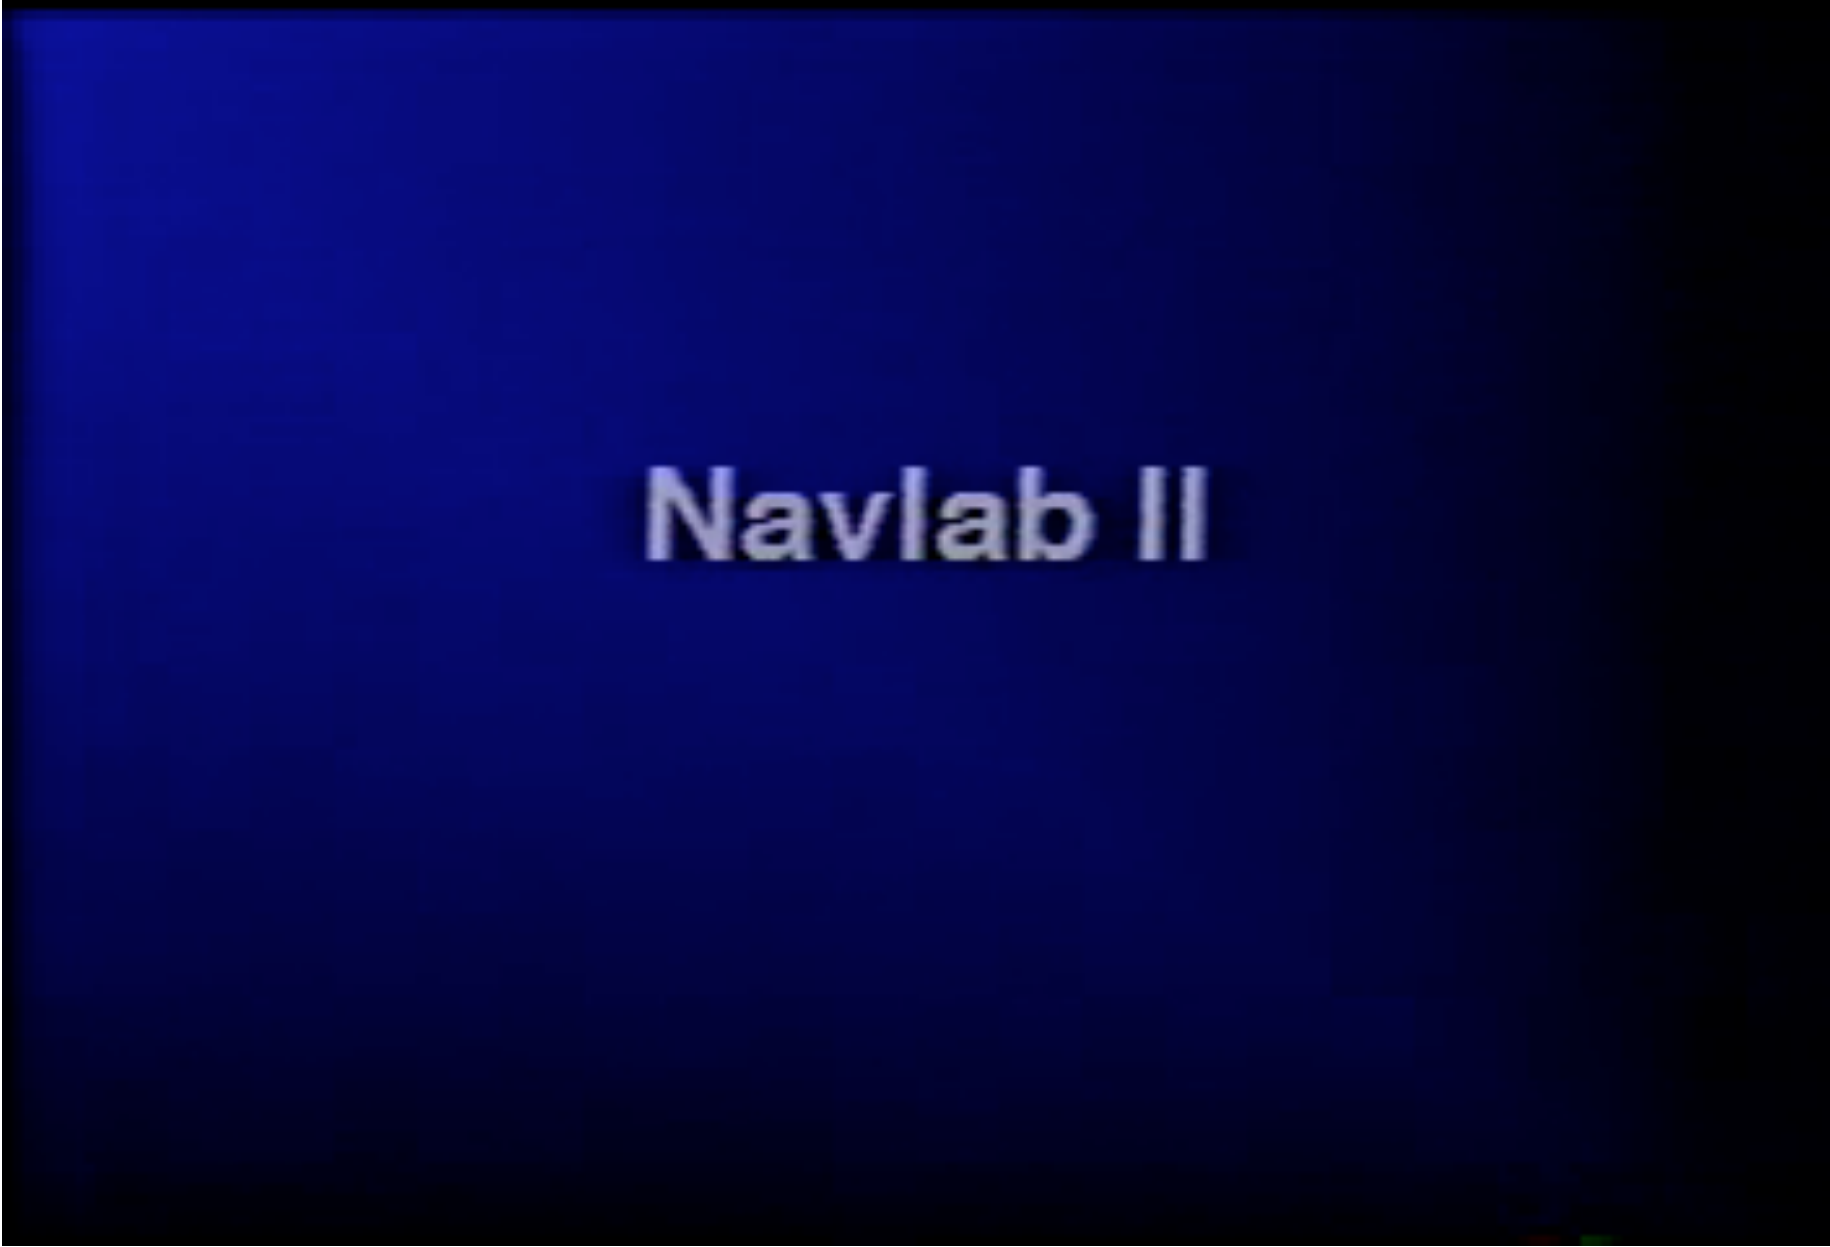
Algorithm summary

- Given: m images and n features \mathbf{x}_{ij}
- For each image i , center the feature coordinates
- Construct a $2m \times n$ measurement matrix \mathbf{D} :
 - Column j contains the projection of point j in all views
 - Row i contains one coordinate of the projections of all the n points in image i
- Factorize \mathbf{D} :
 - Compute SVD: $\mathbf{D} = \mathbf{U} \mathbf{W} \mathbf{V}^T$
 - Create \mathbf{U}_3 by taking the first 3 columns of \mathbf{U}
 - Create \mathbf{V}_3 by taking the first 3 columns of \mathbf{V}
 - Create \mathbf{W}_3 by taking the upper left 3×3 block of \mathbf{W}
- Create the motion and shape matrices:
$$\mathbf{M} = \mathbf{U}_3, \quad \mathbf{S} = \Sigma_3 \mathbf{V}_3^T$$
- Eliminate affine ambiguity

Results



Results




Navlab II

Outline

- Bundle Adjustment
- Ambiguities in Reconstruction
- Affine Factorization
- Extensions

Dealing with missing data


- So far, we have assumed that all points are visible in all views
- In reality, the measurement matrix typically looks something like this:



Sequential structure from motion

- Intuition: exploit low-rank redundancy of matrix
- Possible solution: decompose matrix into dense sub-blocks, factorize each sub-block, and fuse the results
 - Finding dense maximal sub-blocks of the matrix is NP-complete (equivalent to finding maximal cliques in a graph)
- Incremental bilinear refinement

(1) Perform factorization on a dense sub-block




Sequential structure from motion

- Intuition: exploit low-rank redundancy of matrix
- Possible solution: decompose matrix into dense sub-blocks, factorize each sub-block, and fuse the results
 - Finding dense maximal sub-blocks of the matrix is NP-complete (equivalent to finding maximal cliques in a graph)
- Incremental bilinear refinement

- (2) Solve for a new 3D point visible by at least two known cameras (linear least squares)

$$\min_S ||D - MS||^2$$




Sequential structure from motion

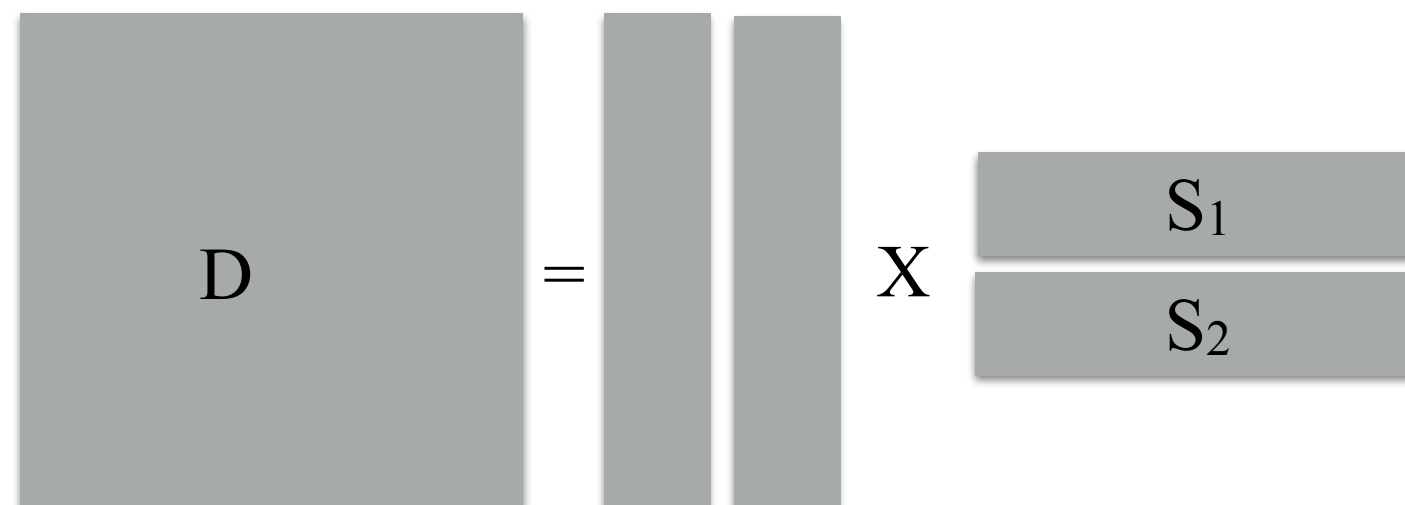
- Intuition: exploit low-rank redundancy of matrix
- Possible solution: decompose matrix into dense sub-blocks, factorize each sub-block, and fuse the results
 - Finding dense maximal sub-blocks of the matrix is NP-complete (equivalent to finding maximal cliques in a graph)
- Incremental bilinear refinement

- (3) Solve for a new camera that sees at least three known 3D points (linear least squares)

$$\min_M ||D - MS||^2$$



Nonrigid structure motion


$$D = \begin{bmatrix} \alpha_1 M & \alpha_2 M \end{bmatrix} X \begin{bmatrix} S_1 \\ S_2 \end{bmatrix}$$

Assume 3D shapes are a linear combination of K basis shapes

$$S = \sum_{k=1}^K \alpha_k S_k$$

For each image, we need to solve for camera matrix *and* scaling coefficients

Simply relax rank constraint on measurement matrix from 3 to $3K$!

Dance Dataset

Mocap data with 75 points, 264 frames, $K = 5$

2D Input Data



Two Views of 3D Reconstruction



● Ground Truth

○ Computed Structure

Projective structure from motion

- Given: m images of n fixed 3D points

$$\mathbf{x}_{ij} \equiv \mathbf{M}_i \mathbf{X}_j, \quad i = 1, \dots, m, \quad j = 1, \dots, n$$

- Problem: estimate m projection matrices \mathbf{M}_i and n 3D points \mathbf{X}_j from the mn correspondences \mathbf{x}_{ij}
- With no calibration info, cameras and points can only be recovered up to a 4x4 projective transformation \mathbf{Q} :

$$\mathbf{X} \rightarrow \mathbf{QX}, \quad \mathbf{M} \rightarrow \mathbf{MQ}^{-1}$$

- We can solve for structure and motion when

$$2mn \geq 11m + 3n - 15$$

- For two cameras, at least 7 points are needed

Projective SFM: Two-camera case


- Compute fundamental matrix F between the two views (from 7 or more correspondences)
- First camera matrix: $M = [I|0]$
- Second camera matrix: $M' = [A_{3 \times 3} | b_{3 \times 1}]$
- Then one can compute A, b from F as follows:
 - b is the right null vector, or epipole ($F^T b = 0$)
 - $A = -\hat{b}F$ (where hat notation refers to skew symmetric matrix)

For proof, see F & P Sec 8.3

For more than 2 images, things get more implicated

Back to bundle adjustment

Minimize reprojection error over multiple 3D points and cameras



$$\min_{\mathbf{X}_1, \mathbf{X}_2, \dots, M_1, M_2, \dots} \sum_{i=1}^m \sum_{j=1}^n \|\mathbf{x}_{ij} - Proj(\mathbf{X}_j, M_i)\|^2$$

Self-calibration

- Self-calibration (auto-calibration) is the process of determining intrinsic camera parameters directly from uncalibrated images
- For example, when the images are acquired by a single moving camera, we can use the constraint that the intrinsic parameter matrix remains fixed for all the images
 - Compute initial projective reconstruction and find 3D projective transformation matrix \mathbf{Q} such that all camera matrices are in the form $\mathbf{M}_i = \mathbf{K} [\mathbf{R}_i | \mathbf{t}_i]$
- Can use constraints on the form of the calibration matrix: orthogonal image axis
- Can use vanishing points

Review: Structure from motion

- Ambiguity
- Affine structure from motion
 - Factorization
- Dealing with missing data
 - Incremental structure from motion
- Projective structure from motion
 - Bundle adjustment
 - Self-calibration

Summary: 3D geometric vision

- Single-view geometry
 - The pinhole camera model
 - Variation: orthographic projection
 - The perspective projection matrix
 - Intrinsic parameters
 - Extrinsic parameters
 - Calibration
- Multiple-view geometry
 - Triangulation
 - The epipolar constraint
 - Essential matrix and fundamental matrix
 - Stereo
 - Binocular, multi-view
 - Structure from motion
 - Reconstruction ambiguity
 - Affine SFM
 - Projective SFM